


Data AI DLC for Basware — Technical Vision

Rapid-build demo for presale use (Basware Discovery Phase 1, "AI DLC Framework and Demo")

1. Purpose

Basware project would benefit from AI acceleration like DLC where part of responsibilities are safely assigned to ClaudeCode or another agentic system. N-iX demo is intended to show an AI-assisted loop that enforces "define before build," with AI absorbing the mechanical work so human effort concentrates on definitions and review:



It's a working demo instance on N-iX accounts, not a production DLC design.

2. Requirements

- Jira ticket picked up and interpreted correctly; agent writes a technical design document *before* code
- Implementation + tests generated and run; PR opened is reviewable by developers
- Steering docs (CLAUDE.md / overview.md / agent instructions) are visible and explainable — this is what maps to Basware's "no shared definitions" pain.
- Jira ticket status updates automatically reflecting the progress
- Runs entirely on N-iX-controlled infra/data, reproducible on demand, buildable in ~1-3 day


3. Constraints


Constraint	Implication
Building the demo from a known reference pattern, not from scratch	Public components (Databricks/Claude skills, Atlassian MCP etc) plus re-authored steering docs go into this build
No real Basware data/definitions yet	Using synthetic data, loosely modeled on invoice-lifecycle entities (customer, contract, invoice)
Basware's actual stack (Jira? GitHub? Azure DevOps?) unconfirmed	Demo runs on N-iX-owned accounts; presented as illustrative, not Basware-specific

4. Approach


Track A — DLC content (architect + framework SME): The framework SME flags which pieces are public so the architect can source equivalents directly, then outlines the steering doc structure (~600+ lines in the reference pattern) from experience. The architect drafts the N-iX/toy-domain version; a peer reviews for coherence.

Track B — Infrastructure (architect):





The loop:



Important: **agents over-infer on incomplete context.** Where steering docs are vague, the agent fills gaps with plausible-but-wrong assumptions — the same failure mode behind the ARR rework, just caught earlier and visibly, in a technical design document.

5. Scope of Work

1. Collect public AI artefacts (Databricks skills, MCP definitions) - architect
2. Reconstruct steering docs - architect and framework SME
3. Databricks trial workspace, GitHub repo, Jira project + demo ticket(s) - architect
4. Databricks↔GitHub and Jira↔Claude Code integrations - architect
5. Local Claude Code setup + synthetic Bronze data - architect
6. One full end-to-end run: technical design document → implementation → tests → PR → Jira update - architect
7. Integrity review of output - framework SME or peer engineer

Out of scope: CI/CD, multi-domain canonical model, real Basware data/logic, production security/governance review.

6. Demo Scenario (10–15 min)

1. Open with the ARR incident.
2. Show a Jira ticket for a deliberately ambiguous metric (*count of active contracts per customer*).
3. Show the steering docs — "if 'customer' isn't defined here, the agent guesses, same as a human."
4. Walk through the generated technical design document (the artifact that would've caught ARR earlier).
5. Walk through the PR (schema/pipeline/tests, reviewable).
6. Show the Jira ticket auto-close.
7. Land the point: this makes "define before build" structural, at AI speed, across all 16 remaining KPIs.

7. Estimates

Task	Effort
Collecting public artefacts	3–4 hrs
Steering doc reconstruction	3–4 hrs
Databricks/GitHub/Jira provisioning	2–3 hrs
Integrations	2–3 hrs
Claude Code setup + toy data	2–3 hrs
Demo ticket + first full run	2–3 hrs

Integrity review	1 hr
Total (parallelized)	Up to 24 hr

8. Risks

- **Steering doc quality** — mitigated by peer review before any external use.
- **Toy-data realism** — offset by naming "customer/contract" explicitly, echoing Basware's own language.
- **Partially known target stack** — demo is tool-illustrative; confirm real tooling during Discovery.
- **First-run steering gaps** — expected; treat as part of the narrative, not something to hide.